

Convolutional Neural Network - Derivations & Proofs

Paul F. Roysdon, Ph.D.

Contents

1 Mathematical Derivations & Proofs	1
1.1 Introduction	1
1.2 Data and Notation	1
1.3 Model Formulation: Discrete Convolution (Cross-Correlation)	2
1.4 Backpropagation Through Convolution	2
1.5 Pooling Layers: Forward and Backward	3
1.6 Classification Head and Loss	4
1.7 Proofs & Identities	4
1.8 Computational Aspects	4
1.9 Algorithm (Supervised CNN Training)	5
1.10 Summary of Variables and Their Dimensions	5
1.11 Summary	5

1 Mathematical Derivations & Proofs

1.1 Introduction

A Convolutional Neural Network (CNN) is a deep architecture for signals on regular grids (e.g., images) built from convolutional layers, elementwise nonlinearities, and pooling/normalization. A convolutional layer applies a bank of learned local, translation-equivariant linear filters to the input, followed by a pointwise nonlinearity. We derive discrete convolution precisely (with padding/stride/dilation), prove translation equivariance, show matrix (Toeplitz/`im2col`) equivalence, and give full backpropagation formulas for weights, biases, and inputs. We also derive pooling forward/backward and a standard classification head with cross-entropy loss. All variables and dimensions are explicit.

1.2 Data and Notation

Let a mini-batch of inputs be

$$\mathbf{X} \in \mathbb{R}^{N \times H \times W \times C_{\text{in}}} \quad (\text{batch, height, width, channels}),$$

with N samples, spatial size (H, W) , and C_{in} channels. A filter bank (kernel) is

$$\mathbf{K} \in \mathbb{R}^{k_h \times k_w \times C_{\text{in}} \times C_{\text{out}}}, \quad \mathbf{b} \in \mathbb{R}^{C_{\text{out}}},$$

and hyperparameters are *padding* $(p_h, p_w) \in \mathbb{Z}_{\geq 0}^2$, *stride* $(s_h, s_w) \in \mathbb{Z}_{\geq 1}^2$, and *dilation* $(d_h, d_w) \in \mathbb{Z}_{\geq 1}^2$. Define the *effective receptive field* per axis

$$r_h = d_h(k_h - 1) + 1, \quad r_w = d_w(k_w - 1) + 1.$$

Then the output spatial size is

$$H_{\text{out}} = \left\lfloor \frac{H + 2p_h - r_h}{s_h} \right\rfloor + 1, \quad W_{\text{out}} = \left\lfloor \frac{W + 2p_w - r_w}{s_w} \right\rfloor + 1,$$

and the layer output tensor is

$$\mathbf{Y} \in \mathbb{R}^{N \times H_{\text{out}} \times W_{\text{out}} \times C_{\text{out}}}.$$

Throughout, bold capitals denote tensors/matrices; indices are zero- or one-based as convenient (we will write integer equalities explicitly).

1.3 Model Formulation: Discrete Convolution (Cross-Correlation)

Deep-learning libraries implement *cross-correlation* (no kernel flip) by default. Define the padded input $\mathbf{X}^{(\text{pad})} \in \mathbb{R}^{N \times (H+2p_h) \times (W+2p_w) \times C_{\text{in}}}$ by zero-padding \mathbf{X} . For $n \in [N]$, $(i, j) \in [H_{\text{out}}] \times [W_{\text{out}}]$, $c \in [C_{\text{out}}]$,

$$\mathbf{Y}[n, i, j, c] = \mathbf{b}[c] + \sum_{u=0}^{k_h-1} \sum_{v=0}^{k_w-1} \sum_{c'=0}^{C_{\text{in}}-1} \mathbf{K}[u, v, c', c] \mathbf{X}^{(\text{pad})}[n, is_h + ud_h, js_w + vd_w, c']. \quad (1)$$

When $d_h = d_w = s_h = s_w = 1$ and $p_h = p_w = 0$, Eqn. (1) reduces to the classical unit-stride valid correlation.

Matrix (im2col/Toeplitz) equivalence. Unfold each output location (i, j) into a row that collects the corresponding receptive-field entries of $\mathbf{X}^{(\text{pad})}$ across channels. This defines a linear operator

$$\mathcal{T}(\mathbf{X}) \in \mathbb{R}^{(NH_{\text{out}}W_{\text{out}}) \times (k_h k_w C_{\text{in}})},$$

and stack kernels column-wise $\text{vec}(\mathbf{K}) \in \mathbb{R}^{(k_h k_w C_{\text{in}}) C_{\text{out}}}$. Then

$$\text{mat}(\mathbf{Y}) = \mathcal{T}(\mathbf{X}) \text{vec}(\mathbf{K}) + \mathbf{1} \mathbf{b}^T, \quad (2)$$

where $\text{mat}(\mathbf{Y}) \in \mathbb{R}^{(NH_{\text{out}}W_{\text{out}}) \times C_{\text{out}}}$ flattens (n, i, j) into rows and $\mathbf{1}$ is the all-ones column. Thus convolution is a *linear map* w.r.t. both \mathbf{X} and \mathbf{K} .

Translation equivariance (stride = 1, no padding-edge effects). Let τ_Δ shift \mathbf{X} by integer offset $\Delta = (\Delta_h, \Delta_w)$, i.e., $(\tau_\Delta \mathbf{X})[n, a, b, c] = \mathbf{X}[n, a - \Delta_h, b - \Delta_w, c]$ where defined, else 0. Then for p large enough to avoid border clipping,

$$\text{Conv}(\tau_\Delta \mathbf{X}; \mathbf{K}, \mathbf{b})[n, i, j, c] = \tau_\Delta(\text{Conv}(\mathbf{X}; \mathbf{K}, \mathbf{b}))[n, i, j, c], \quad (3)$$

i.e., shifts commute with convolution: the layer is translation *equivariant*.

Proof. Substitute $\tau_\Delta \mathbf{X}$ into Eqn. (1) and re-index the sums: the output indices shift by the same Δ because the kernel support is finite and shared over all locations. \blacksquare

Nonlinearity and block. A standard block is

$$\mathbf{Z} = \sigma(\text{Conv}(\mathbf{X}; \mathbf{K}, \mathbf{b})), \quad \text{e.g., } \sigma = \text{ReLU} : \sigma(t) = \max\{t, 0\} \text{ (elementwise)}.$$

1.4 Backpropagation Through Convolution

Let \mathcal{L} be the scalar loss. Define the *output sensitivity*

$$\mathbf{G}_Y \triangleq \frac{\partial \mathcal{L}}{\partial \mathbf{Y}} \in \mathbb{R}^{N \times H_{\text{out}} \times W_{\text{out}} \times C_{\text{out}}}.$$

All gradients follow from linearity (plus the chain rule for nonlinearities).

Bias gradient. By Eqn. (1), $\partial \mathbf{Y}[n, i, j, c] / \partial \mathbf{b}[c'] = \mathbf{1}\{c = c'\}$, hence

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}[c]} = \sum_{n, i, j} \mathbf{G}_Y[n, i, j, c]. \quad (4)$$

Kernel (weight) gradient. From Eqn. (1), $\partial \mathbf{Y}[n, i, j, c] / \partial \mathbf{K}[u, v, c', c] = \mathbf{X}^{(\text{pad})}[n, is_h + ud_h, js_w + vd_w, c']$. Therefore

$$\frac{\partial \mathcal{L}}{\partial \mathbf{K}[u, v, c', c]} = \sum_{n=0}^{N-1} \sum_{i=0}^{H_{\text{out}}-1} \sum_{j=0}^{W_{\text{out}}-1} \mathbf{G}_Y[n, i, j, c] \mathbf{X}^{(\text{pad})}[n, is_h + ud_h, js_w + vd_w, c']. \quad (5)$$

Interpretation. $\partial \mathcal{L} / \partial \mathbf{K}$ is a (multi-channel) *cross-correlation* between $\mathbf{X}^{(\text{pad})}$ and \mathbf{G}_Y , accumulated over batch and output sites.

Input gradient (transposed convolution). For a fixed input position (a, b, c') (in the *padded* tensor),

$$\frac{\partial \mathbf{Y}[n, i, j, c]}{\partial \mathbf{X}^{(\text{pad})}[n, a, b, c']} = \sum_{u, v} \mathbf{K}[u, v, c', c] \mathbf{1}\{a = is_h + ud_h, b = js_w + vd_w\}.$$

Thus

$$\frac{\partial \mathcal{L}}{\partial \mathbf{X}^{(\text{pad})}[n, a, b, c']} = \sum_{c=0}^{C_{\text{out}}-1} \sum_{u=0}^{k_h-1} \sum_{v=0}^{k_w-1} \mathbf{K}[u, v, c', c] \mathbf{G}_Y\left[n, \frac{a - ud_h}{s_h}, \frac{b - vd_w}{s_w}, c\right], \quad (6)$$

where terms with non-integer indices are zero (Kronecker constraints). After computing $\partial \mathcal{L} / \partial \mathbf{X}^{(\text{pad})}$, crop the padding to obtain $\partial \mathcal{L} / \partial \mathbf{X}$.

Compact forms. *Proof.* Equations Eqn. (5) and Eqn. (6) can be implemented as standard convolutions:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{K}} = \text{XCorr}_{n, i, j}(\mathbf{X}^{(\text{pad})}, \mathbf{G}_Y), \quad (7)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{X}} = \text{ConvTransp}(\mathbf{G}_Y, \mathbf{K}; s, p, d), \quad (\text{transpose of forward w.r.t. } \mathbf{X}). \quad (8)$$

Alternatively, via Eqn. (2):

$$\text{vec}\left(\frac{\partial \mathcal{L}}{\partial \mathbf{K}}\right) = \mathcal{T}(\mathbf{X})^\top \text{vec}(\mathbf{G}_Y), \quad \frac{\partial \mathcal{L}}{\partial \mathbf{b}} = (\text{mat}(\mathbf{G}_Y)^\top \mathbf{1}), \quad (9)$$

$$\text{mat}\left(\frac{\partial \mathcal{L}}{\partial \mathbf{X}}\right) = \mathcal{T}^\top(\text{vec}(\mathbf{K})) \text{mat}(\mathbf{G}_Y), \quad (10)$$

where \mathcal{T}^\top denotes the adjoint (col2im) operator. These equalities follow from standard matrix calculus for $\mathcal{L} = \phi(\mathcal{T}(\mathbf{X})\theta + \mathbf{1}\mathbf{b}^\top)$. \blacksquare

Nonlinearity backprop. For $\mathbf{Z} = \sigma(\mathbf{Y})$ and upstream $\mathbf{G}_Z = \partial \mathcal{L} / \partial \mathbf{Z}$,

$$\mathbf{G}_Y = \mathbf{G}_Z \odot \sigma'(\mathbf{Y}).$$

For ReLU, $\sigma'(t) = \mathbf{1}\{t > 0\}$.

1.5 Pooling Layers: Forward and Backward

Let $\text{Pool}_{k_h \times k_w}^{(s_h, s_w)}$ operate per-channel on non-overlapping or strided windows.

Average pooling. Forward:

$$\mathbf{Y}[n, i, j, c] = \frac{1}{k_h k_w} \sum_{u=0}^{k_h-1} \sum_{v=0}^{k_w-1} \mathbf{X}[n, is_h + u, js_w + v, c].$$

Backward distributes gradients uniformly:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{X}[n, a, b, c]} = \sum_{i,j} \frac{1}{k_h k_w} \mathbf{G}_Y[n, i, j, c] \mathbf{1}\{a \in [is_h, is_h+k_h-1], b \in [js_w, js_w+k_w-1]\}.$$

Max pooling. Forward selects maxima per window; backward routes gradients to argmax locations:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{X}[n, a, b, c]} = \sum_{i,j} \mathbf{G}_Y[n, i, j, c] \mathbf{1}\left\{\mathbf{X}[n, a, b, c] = \max_{u,v} \mathbf{X}[n, is_h+u, js_w+v, c]\right\}.$$

1.6 Classification Head and Loss

Let the final convolutional block output $\mathbf{H} \in \mathbb{R}^{N \times H' \times W' \times C'}$. **Global average pooling** (GAP) aggregates spatially:

$$\mathbf{g}[n, c] = \frac{1}{H'W'} \sum_{i=0}^{H'-1} \sum_{j=0}^{W'-1} \mathbf{H}[n, i, j, c], \quad \mathbf{g} \in \mathbb{R}^{N \times C'}.$$

A linear readout to K classes:

$$\mathbf{r} = \mathbf{g} \mathbf{W}_{\text{cls}} + \mathbf{b}_{\text{cls}}, \quad \mathbf{W}_{\text{cls}} \in \mathbb{R}^{C' \times K}, \quad \mathbf{b}_{\text{cls}} \in \mathbb{R}^K.$$

With one-hot labels $\mathbf{y} \in \{0, 1\}^{N \times K}$ and softmax $p_{nk} = \exp(r_{nk}) / \sum_t \exp(r_{nt})$, the cross-entropy loss is

$$\mathcal{L} = - \sum_{n=1}^N \sum_{k=1}^K \mathbf{y}_{nk} \log p_{nk}, \quad \Rightarrow \quad \frac{\partial \mathcal{L}}{\partial \mathbf{r}} = \mathbf{p} - \mathbf{y}.$$

Gradients propagate through GAP by uniform redistribution over spatial sites.

1.7 Proofs & Identities

Convolution is a linear map. *Proof.* For fixed (\mathbf{K}, \mathbf{b}) , Eqn. (1) is affine in \mathbf{X} : $\text{Conv}(\alpha \mathbf{X}_1 + \beta \mathbf{X}_2) = \alpha \text{Conv}(\mathbf{X}_1) + \beta \text{Conv}(\mathbf{X}_2)$ plus \mathbf{b} . Similarly, for fixed \mathbf{X} , it is linear in \mathbf{K} . ■

Translation equivariance (formal). *Proof.* Let $s_h = s_w = 1$ and padding large enough to avoid boundary truncation. For any Δ , by substituting $\tau_\Delta \mathbf{X}$ into Eqn. (1) and re-indexing $i' = i + \Delta_h$, $j' = j + \Delta_w$, $\mathbf{Y}_{\tau_\Delta \mathbf{X}}[i, j, c] = \mathbf{Y}_{\mathbf{X}}[i - \Delta_h, j - \Delta_w, c]$, which is the statement of Eqn. (3). ■

Backprop correctness (matrix view). *Proof.* Let $\mathbf{Y} = \mathcal{T}(\mathbf{X}) \Theta + \mathbf{1} \mathbf{b}^\top$ with $\Theta = \text{vec}(\mathbf{K})$. For any scalar loss $\mathcal{L}(\mathbf{Y})$, matrix calculus gives $\partial \mathcal{L} / \partial \Theta = \mathcal{T}(\mathbf{X})^\top \text{vec}(\partial \mathcal{L} / \partial \mathbf{Y})$ and $\partial \mathcal{L} / \partial \mathbf{X} = \mathcal{T}^*(\Theta (\partial \mathcal{L} / \partial \mathbf{Y})^\top)$, where \mathcal{T}^* is the adjoint operator (col2im). These match Eqn. (9). ■

1.8 Computational Aspects

A direct convolution costs

$$\mathcal{O}(N H_{\text{out}} W_{\text{out}} k_h k_w C_{\text{in}} C_{\text{out}}).$$

Implementations use im2col+GEMM, Winograd (for small 3×3), or FFT for large kernels. Memory/computation scale linearly with $C_{\text{in}} C_{\text{out}}$ and kernel area.

1.9 Algorithm (Supervised CNN Training)

1. **Input:** batch (\mathbf{X}, \mathbf{y}) , parameters $\{\mathbf{K}^{(\ell)}, \mathbf{b}^{(\ell)}\}_{\ell=1}^L$ and head $(\mathbf{W}_{\text{cls}}, \mathbf{b}_{\text{cls}})$; hyperparameters (p, s, d) , learning rate η .
2. **Forward:** For layers $\ell = 1, \dots, L$,

$$\mathbf{Y}^{(\ell)} \leftarrow \text{Conv}(\mathbf{X}^{(\ell-1)}; \mathbf{K}^{(\ell)}, \mathbf{b}^{(\ell)}), \quad \mathbf{X}^{(\ell)} \leftarrow \sigma(\mathbf{Y}^{(\ell)}),$$

optionally with pooling/normalization between layers. Apply GAP and linear head, then softmax and cross-entropy.

3. **Backward:** Compute $\partial \mathcal{L}/\partial \mathbf{r} = \mathbf{p} - \mathbf{y}$, backprop through head and GAP, then per layer

$$\mathbf{G}_Y^{(\ell)} \leftarrow \mathbf{G}_X^{(\ell)} \odot \sigma'(\mathbf{Y}^{(\ell)}), \quad \frac{\partial \mathcal{L}}{\partial \mathbf{K}^{(\ell)}} \leftarrow \text{Eqn. (5)}, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{b}^{(\ell)}} \leftarrow \text{Eqn. (4)}, \quad \mathbf{G}_X^{(\ell-1)} \leftarrow \text{Eqn. (6) (cropped)}.$$

4. **Update:** Apply an optimizer (SGD/Adam/AdamW) to all parameters.

1.10 Summary of Variables and Their Dimensions

- $\mathbf{X} \in \mathbb{R}^{N \times H \times W \times C_{\text{in}}}$: input batch.
- $\mathbf{K} \in \mathbb{R}^{k_h \times k_w \times C_{\text{in}} \times C_{\text{out}}}$: convolution kernels.
- $\mathbf{b} \in \mathbb{R}^{C_{\text{out}}}$: per-output-channel bias.
- $(p_h, p_w), (s_h, s_w), (d_h, d_w)$: padding, stride, dilation (scalars per axis).
- $\mathbf{Y} \in \mathbb{R}^{N \times H_{\text{out}} \times W_{\text{out}} \times C_{\text{out}}}$: pre-activation outputs.
- $\mathbf{Z} = \sigma(\mathbf{Y})$: post-activation outputs (same shape as \mathbf{Y}).
- $\mathcal{T}(\mathbf{X}) \in \mathbb{R}^{(NH_{\text{out}}W_{\text{out}}) \times (k_h k_w C_{\text{in}})}$: im2col operator.
- $\mathbf{G}_Y = \partial \mathcal{L}/\partial \mathbf{Y}, \mathbf{G}_X = \partial \mathcal{L}/\partial \mathbf{X}$: sensitivities.
- Head: $\mathbf{W}_{\text{cls}} \in \mathbb{R}^{C' \times K}, \mathbf{b}_{\text{cls}} \in \mathbb{R}^K$; logits \mathbf{r} , probabilities \mathbf{p} .

1.11 Summary

From first principles, a convolutional layer is a translation-equivariant, local linear operator with shared weights, expressed by Eqn. (1) (or equivalently the Toeplitz product Eqn. (2)). Backpropagation yields closed-form gradients: the bias sum Eqn. (4), kernel cross-correlation Eqn. (5), and input gradient as a transposed convolution Eqn. (6)/Eqn. (7). Pooling layers are simple linear (average) or selection (max) operators with corresponding adjoints. Stacking conv–nonlinearity–pooling blocks with a classification head and training via cross-entropy completes a CNN, with all dimensions and derivations explicitly specified.