

Graph Convolutional Neural Network - Derivations & Proofs

Paul F. Roysdon, Ph.D.

Contents

1 Mathematical Derivations & Proofs	1
1.1 Introduction	1
1.2 Data and Notation	1
1.3 Model Formulation: Graph Convolution from Spectral First Principles	2
1.4 Permutation Equivariance (Graph Isomorphism Invariance)	3
1.5 Message-Passing View and k -Hop Locality	3
1.6 Training Objective and Semi-Supervised Setting	3
1.7 Backpropagation Through a GCN Layer	3
1.8 Computational Aspects	4
1.9 Algorithm (Semi-Supervised Node Classification with GCN)	4
1.10 Proofs and Identities	4
1.11 Extensions (Brief)	5
1.12 Summary of Variables and Their Dimensions	5
1.13 Summary	5

1 Mathematical Derivations & Proofs

1.1 Introduction

A Graph Convolutional Neural Network (GCN) generalizes convolution to signals defined on the vertices of a graph. Starting from the graph Laplacian and the Graph Fourier Transform, one derives spectral graph filters and then obtains a localized, linear-time approximation (Chebyshev polynomials). Specializing to a first-order approximation and a renormalized adjacency produces the widely-used propagation rule

$$\mathbf{H}^{(\ell+1)} = \sigma(\hat{\mathbf{A}} \mathbf{H}^{(\ell)} \mathbf{W}^{(\ell)} + \mathbf{1} (\mathbf{b}^{(\ell)})^\top), \quad \hat{\mathbf{A}} = \tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2}, \quad \tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I},$$

which is *permutation equivariant* and implements *neighborhood averaging* (Laplacian smoothing) followed by a learnable feature mixing. We derive these results, prove locality and permutation equivariance, and provide full backpropagation formulas with all variables and dimensions explicit.

1.2 Data and Notation

Let $G = (V, E)$ be an undirected (possibly weighted) graph with $|V| = n$ nodes. Define:

- **Adjacency:** $\mathbf{A} \in \mathbb{R}^{n \times n}$, where $\mathbf{A}_{ij} \geq 0$ is the edge weight (zero if no edge). For self-loops we will use $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$.
- **Degree:** $\mathbf{D} \in \mathbb{R}^{n \times n}$ diagonal, $\mathbf{D}_{ii} = \sum_j \mathbf{A}_{ij}$; likewise $\tilde{\mathbf{D}}_{ii} = \sum_j \tilde{\mathbf{A}}_{ij}$.
- **(Normalized) Laplacian:** $\mathbf{L} = \mathbf{D} - \mathbf{A}$ and $\mathbf{L}_{\text{sym}} = \mathbf{I} - \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$.
- **Node features:** $\mathbf{X} \in \mathbb{R}^{n \times F_{\text{in}}}$, with F_{in} input feature channels per node (rows index nodes).

- **Layer representations:** $\mathbf{H}^{(0)} \equiv \mathbf{X}$, $\mathbf{H}^{(\ell)} \in \mathbb{R}^{n \times F_\ell}$, $\ell = 0, \dots, L$.
- **Trainable weights/bias:** $\mathbf{W}^{(\ell)} \in \mathbb{R}^{F_\ell \times F_{\ell+1}}$, $\mathbf{b}^{(\ell)} \in \mathbb{R}^{F_{\ell+1}}$.

We use $\mathbf{1} \in \mathbb{R}^{n \times 1}$ for the all-ones column. Nonlinearities $\sigma(\cdot)$ act elementwise.

1.3 Model Formulation: Graph Convolution from Spectral First Principles

Graph Fourier Transform (GFT). Since \mathbf{L}_{sym} is real symmetric, it has an eigendecomposition

$$\mathbf{L}_{\text{sym}} = \mathbf{U} \Lambda \mathbf{U}^\top, \quad \mathbf{U} \in \mathbb{R}^{n \times n} \text{ orthonormal}, \quad \Lambda = \text{diag}(\lambda_1, \dots, \lambda_n), \quad \lambda_k \in [0, 2].$$

For a graph signal $\mathbf{f} \in \mathbb{R}^n$, the GFT is $\hat{\mathbf{f}} = \mathbf{U}^\top \mathbf{f}$ and inverse is $\mathbf{f} = \mathbf{U} \hat{\mathbf{f}}$.

Spectral graph filtering. A *spectral filter* with transfer function $g(\cdot)$ acts by

$$g \star \mathbf{f} = \mathbf{U} g(\Lambda) \mathbf{U}^\top \mathbf{f}, \quad g(\Lambda) = \text{diag}(g(\lambda_1), \dots, g(\lambda_n)). \quad (1)$$

For matrix-valued features $\mathbf{H} \in \mathbb{R}^{n \times F}$, apply Eqn. (1) columnwise.

Polynomial filters \Rightarrow spatial locality. Let $g(\lambda) = \sum_{k=0}^K \theta_k T_k(\tilde{\lambda})$ be a degree- K polynomial in the (rescaled) eigenvalue $\tilde{\lambda} \in [-1, 1]$ with Chebyshev polynomials T_k .¹ Then

$$g(\tilde{\mathbf{L}}) \mathbf{H} = \sum_{k=0}^K \theta_k T_k(\tilde{\mathbf{L}}) \mathbf{H}, \quad (2)$$

and $T_k(\tilde{\mathbf{L}})$ is a k -hop localized operator:

$$[T_k(\tilde{\mathbf{L}})]_{ij} = 0 \quad \text{whenever the graph distance } \text{dist}(i, j) > k.$$

Proof. $T_0(\tilde{\mathbf{L}}) = \mathbf{I}$ and $T_1(\tilde{\mathbf{L}}) = \tilde{\mathbf{L}}$. The recurrence $T_{k+1}(\tilde{\mathbf{L}}) = 2\tilde{\mathbf{L}}T_k(\tilde{\mathbf{L}}) - T_{k-1}(\tilde{\mathbf{L}})$ yields a polynomial in $\tilde{\mathbf{L}}$ of total degree $k+1$. As $\tilde{\mathbf{L}}$ is a sparsified version of \mathbf{L}_{sym} with nonzeros only on edges and self-loops, any degree- k polynomial connects at most k hops. ■

First-order (K=1) approximation \Rightarrow GCN propagation. Take $K = 1$ and expand $g(\tilde{\mathbf{L}}) \approx \theta_0 T_0(\tilde{\mathbf{L}}) + \theta_1 T_1(\tilde{\mathbf{L}}) = \theta_0 \mathbf{I} + \theta_1 \tilde{\mathbf{L}}$. Undoing the rescaling gives (up to constants)

$$g(\mathbf{L}_{\text{sym}}) \mathbf{H} \approx \theta_0 \mathbf{I} \mathbf{H} - \theta_1 \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} \mathbf{H}. \quad (3)$$

Setting $\theta = \theta_0 = \theta_1$ and adding self-loops $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ (so that the identity term is absorbed into $\tilde{\mathbf{A}}$), one obtains the *renormalized* operator

$$\hat{\mathbf{A}} \triangleq \tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2}, \quad \tilde{\mathbf{D}}_{ii} = \sum_j \tilde{\mathbf{A}}_{ij}. \quad (4)$$

Finally, mix feature channels with a trainable matrix \mathbf{W} and bias \mathbf{b} :

$$\mathbf{Z}^{(\ell)} = \hat{\mathbf{A}} \mathbf{H}^{(\ell)} \mathbf{W}^{(\ell)} + \mathbf{1} (\mathbf{b}^{(\ell)})^\top, \quad \mathbf{H}^{(\ell+1)} = \sigma(\mathbf{Z}^{(\ell)}).$$

(5)

Thus a GCN layer performs *neighborhood averaging* (via $\hat{\mathbf{A}}$) followed by a per-node linear map.

¹Use the rescaled Laplacian $\tilde{\mathbf{L}} = \frac{2}{\lambda_{\max}} \mathbf{L}_{\text{sym}} - \mathbf{I}$ so that its spectrum lies in $[-1, 1]$.

Interpretation (Laplacian smoothing). For one feature channel ($F_\ell = 1$) and $\sigma = \text{id}$, $\hat{\mathbf{A}}\mathbf{h}$ is a degree-normalized average of each node's own value and its neighbors', which is a one-step smoothing that reduces the Dirichlet energy $\mathbf{h}^\top \mathbf{L}_{\text{sym}} \mathbf{h}$.

1.4 Permutation Equivariance (Graph Isomorphism Invariance)

Let $\mathbf{P} \in \mathbb{R}^{n \times n}$ be a permutation matrix (relabeling of nodes). Under relabeling,

$$\mathbf{A}' = \mathbf{P}\mathbf{A}\mathbf{P}^\top, \quad \tilde{\mathbf{A}}' = \mathbf{P}\tilde{\mathbf{A}}\mathbf{P}^\top, \quad \tilde{\mathbf{D}}' = \mathbf{P}\tilde{\mathbf{D}}\mathbf{P}^\top, \quad \hat{\mathbf{A}}' = \mathbf{P}\hat{\mathbf{A}}\mathbf{P}^\top, \quad \mathbf{H}'^{(\ell)} = \mathbf{P}\mathbf{H}^{(\ell)}.$$

Claim. The GCN layer Eqn. (5) is permutation equivariant:

$$\mathbf{H}'^{(\ell+1)} = \mathbf{P}\mathbf{H}^{(\ell+1)}.$$

Proof.

$$\mathbf{Z}'^{(\ell)} = \hat{\mathbf{A}}'\mathbf{H}'^{(\ell)}\mathbf{W}^{(\ell)} + \mathbf{1}(\mathbf{b}^{(\ell)})^\top = \mathbf{P}\hat{\mathbf{A}}\mathbf{P}^\top\mathbf{P}\mathbf{H}^{(\ell)}\mathbf{W}^{(\ell)} + \mathbf{1}(\mathbf{b}^{(\ell)})^\top = \hat{\mathbf{A}}\mathbf{H}^{(\ell)}\mathbf{W}^{(\ell)} + \mathbf{1}(\mathbf{b}^{(\ell)})^\top.$$

Since $\mathbf{1}$ is invariant to permutation ($\mathbf{P}^\top \mathbf{1} = \mathbf{1}$), we can write $\mathbf{1}(\mathbf{b}^\top) = \mathbf{P}\mathbf{1}(\mathbf{b}^\top)$; hence $\mathbf{Z}'^{(\ell)} = \mathbf{P}\mathbf{Z}^{(\ell)}$ and by elementwise σ , $\mathbf{H}'^{(\ell+1)} = \mathbf{P}\mathbf{H}^{(\ell+1)}$. \blacksquare

1.5 Message-Passing View and k -Hop Locality

Equation (5) can be written nodewise:

$$\mathbf{z}_i^{(\ell)} = \sum_{j \in \mathcal{N}(i) \cup \{i\}} \frac{1}{\sqrt{\tilde{d}_i \tilde{d}_j}} \mathbf{h}_j^{(\ell)} \mathbf{W}^{(\ell)} + \mathbf{b}^{(\ell)}, \quad \mathbf{h}_i^{(\ell+1)} = \sigma(\mathbf{z}_i^{(\ell)}), \quad (6)$$

where $\tilde{d}_i = \tilde{\mathbf{D}}_{ii}$ and $\mathcal{N}(i)$ are neighbors of i . Thus a single layer aggregates *1-hop* messages; composing L layers yields L -hop receptive fields. More generally, polynomial filters of degree K (Eqn. (2)) are K -hop localized.

1.6 Training Objective and Semi-Supervised Setting

Let $\mathcal{S} \subseteq \{1, \dots, n\}$ be the index set of labeled nodes, K classes, and final layer produce logits $\mathbf{R} \in \mathbb{R}^{n \times K}$ with row-softmax \mathbf{P} :

$$\mathbf{R} = \mathbf{H}^{(L)}\mathbf{W}^{(L)} + \mathbf{1}(\mathbf{b}^{(L)})^\top, \quad \mathbf{P}_{ik} = \frac{\exp(\mathbf{R}_{ik})}{\sum_{t=1}^K \exp(\mathbf{R}_{it})}.$$

Using one-hot targets $\mathbf{Y} \in \{0, 1\}^{n \times K}$, the cross-entropy loss restricted to \mathcal{S} is

$$\mathcal{L} = - \sum_{i \in \mathcal{S}} \sum_{k=1}^K \mathbf{Y}_{ik} \log \mathbf{P}_{ik} + \sum_{\ell=0}^L \frac{\lambda_\ell}{2} \|\mathbf{W}^{(\ell)}\|_F^2 \quad (\text{optional } \ell_2 \text{ regularization}). \quad (7)$$

1.7 Backpropagation Through a GCN Layer

Consider one layer (drop superscript ℓ) with pre-activation

$$\mathbf{Z} = \hat{\mathbf{A}}\mathbf{H}\mathbf{W} + \mathbf{1}\mathbf{b}^\top, \quad \mathbf{H}_{\text{next}} = \sigma(\mathbf{Z}), \quad (8)$$

and let $\mathbf{G}_{\text{next}} = \partial \mathcal{L} / \partial \mathbf{H}_{\text{next}} \in \mathbb{R}^{n \times F_{\text{out}}}$. Define $\mathbf{G}_Z = \mathbf{G}_{\text{next}} \odot \sigma'(\mathbf{Z})$ (Hadamard product).

Bias gradient.

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}} = (\mathbf{G}_Z)^\top \mathbf{1} = \sum_{i=1}^n \mathbf{G}_Z(i,:) \in \mathbb{R}^{F_{\text{out}}}.$$
 (9)

Weight gradient. Using $\mathbf{Z} = (\hat{\mathbf{A}}\mathbf{H})\mathbf{W} + \mathbf{1}\mathbf{b}^\top$,

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}} = (\hat{\mathbf{A}}\mathbf{H})^\top \mathbf{G}_Z \in \mathbb{R}^{F_{\text{in}} \times F_{\text{out}}}.$$
 (10)

If $\hat{\mathbf{A}}$ is symmetric (as in Eqn. (4)), we may also write $\frac{\partial \mathcal{L}}{\partial \mathbf{W}} = \mathbf{H}^\top \hat{\mathbf{A}} \mathbf{G}_Z$.

Input gradient.

$$\frac{\partial \mathcal{L}}{\partial \mathbf{H}} = \hat{\mathbf{A}}^\top \mathbf{G}_Z \mathbf{W}^\top = \hat{\mathbf{A}} \mathbf{G}_Z \mathbf{W}^\top \in \mathbb{R}^{n \times F_{\text{in}}}.$$
 (11)

(Optional) Gradient w.r.t. edge-weights. If $\tilde{\mathbf{A}}$ (or $\hat{\mathbf{A}}$) is learnable, use $\frac{\partial \mathcal{L}}{\partial \tilde{\mathbf{A}}} = \mathbf{G}_Z (\mathbf{H}\mathbf{W})^\top$ and chain-rule through $\hat{\mathbf{A}} = \tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2}$ (note: this requires care with the degree normalization).

1.8 Computational Aspects

With sparse $\hat{\mathbf{A}}$, one forward layer costs

$$\mathcal{O}(|E| F_{\text{in}} + n F_{\text{in}} F_{\text{out}})$$

for computing $\hat{\mathbf{A}}\mathbf{H}$ (sparse-dense) and the dense mixing by \mathbf{W} . Training time/space scale linearly in $|E|$ for fixed feature widths. Mini-batching on large graphs typically uses neighborhood sampling (e.g., k -hop subgraphs) to bound computation.

1.9 Algorithm (Semi-Supervised Node Classification with GCN)

1. **Input:** graph (\mathbf{A}, \mathbf{X}) , labeled-node set \mathcal{S} with one-hot labels $\mathbf{Y}[\mathcal{S}, :]$, depth L , hidden widths $\{F_\ell\}$, nonlinearity σ , learning rate η .
2. **Precompute:** $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$, $\tilde{\mathbf{D}}_{ii} = \sum_j \tilde{\mathbf{A}}_{ij}$, $\hat{\mathbf{A}} = \tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2}$ (sparse).
3. **Initialize:** $\mathbf{H}^{(0)} \leftarrow \mathbf{X}$, parameters $\{\mathbf{W}^{(\ell)}, \mathbf{b}^{(\ell)}\}$.
4. **For** epoch = 1, … until convergence:
 - (a) **Forward:** For $\ell = 0, \dots, L-1$, compute $\mathbf{Z}^{(\ell)} = \hat{\mathbf{A}}\mathbf{H}^{(\ell)}\mathbf{W}^{(\ell)} + \mathbf{1}(\mathbf{b}^{(\ell)})^\top$, $\mathbf{H}^{(\ell+1)} = \sigma(\mathbf{Z}^{(\ell)})$. For the final layer, apply softmax to logits.
 - (b) **Loss:** Evaluate \mathcal{L} in Eqn. (7) on \mathcal{S} .
 - (c) **Backward:** Backpropagate using Eqns. (9)–(11).
 - (d) **Update:** Apply an optimizer (SGD/Adam) to all $\{\mathbf{W}^{(\ell)}, \mathbf{b}^{(\ell)}\}$.

1.10 Proofs and Identities

From spectral to spatial (locality). *Proof.* A degree- K polynomial $g(\mathbf{L}_{\text{sym}}) = \sum_{k=0}^K \theta_k \mathbf{L}_{\text{sym}}^k$ only mixes features along paths of length $\leq K$; therefore $g(\mathbf{L}_{\text{sym}})\mathbf{H}$ is K -hop localized. Combining with feature mixing \mathbf{W} preserves locality. ■

Renormalization trick. *Proof.* Starting from Eqn. (3), absorbing the identity with self-loops $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ and re-normalizing to $\hat{\mathbf{A}}$ in Eqn. (4) stabilizes the spectrum of the propagator ($\|\hat{\mathbf{A}}\|_2 \leq 1$), mitigating exploding/vanishing over multiple layers. ■

GCN as neighborhood averaging + linear map. *Proof.* Equation (6) shows $\hat{\mathbf{A}}$ performs a symmetric, degree-weighted average of neighbor features, which is exactly Laplacian smoothing. The learnable \mathbf{W} then recombines channels; σ adds nonlinearity. ■

1.11 Extensions (Brief)

- **Chebyshev GCN (ChebNet).** Use Eqn. (2) with $K > 1$ to obtain K -hop localized filters without eigen-decomposition; compute $T_k(\tilde{\mathbf{L}})\mathbf{H}$ via the three-term recurrence.
- **Edge weights/directions.** For weighted graphs, \mathbf{A} carries weights; for directed graphs, one may use symmetrization or separate in/out normalizations.
- **Dropout and residuals.** Apply dropout to $\mathbf{H}^{(\ell)}$ or edges (DropEdge) and add residual/skip connections to alleviate over-smoothing.

1.12 Summary of Variables and Their Dimensions

- $\mathbf{A} \in \mathbb{R}^{n \times n}$: adjacency (symmetric, nonnegative entries). $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ includes self-loops.
- $\mathbf{D}, \tilde{\mathbf{D}} \in \mathbb{R}^{n \times n}$: degree diagonals; $\tilde{\mathbf{D}}_{ii} = \sum_j \tilde{\mathbf{A}}_{ij}$.
- $\hat{\mathbf{A}} = \tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \in \mathbb{R}^{n \times n}$: renormalized propagator.
- $\mathbf{L}_{\text{sym}} = \mathbf{I} - \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2} \in \mathbb{R}^{n \times n}$: normalized Laplacian.
- $\mathbf{X} = \mathbf{H}^{(0)} \in \mathbb{R}^{n \times F_0}$: input features; $\mathbf{H}^{(\ell)} \in \mathbb{R}^{n \times F_\ell}$ hidden representations.
- $\mathbf{W}^{(\ell)} \in \mathbb{R}^{F_\ell \times F_{\ell+1}}$, $\mathbf{b}^{(\ell)} \in \mathbb{R}^{F_{\ell+1}}$: layer parameters.
- $\mathbf{Z}^{(\ell)} \in \mathbb{R}^{n \times F_{\ell+1}}$: pre-activations; $\mathbf{H}^{(\ell+1)} = \sigma(\mathbf{Z}^{(\ell)})$.
- $\mathbf{R} \in \mathbb{R}^{n \times K}$: final logits; $\mathbf{P} \in \mathbb{R}^{n \times K}$: softmax probabilities.
- $\mathcal{S} \subseteq \{1, \dots, n\}$: labeled-node index set.
- Gradients: $\partial \mathcal{L} / \partial \mathbf{b} \in \mathbb{R}^{F_{\text{out}}}$ (Eqn. (9)), $\partial \mathcal{L} / \partial \mathbf{W} \in \mathbb{R}^{F_{\text{in}} \times F_{\text{out}}}$ (Eqn. (10)), $\partial \mathcal{L} / \partial \mathbf{H} \in \mathbb{R}^{n \times F_{\text{in}}}$ (Eqn. (11)).

1.13 Summary

From first principles: define convolution on graphs via the Laplacian eigenbasis and spectral filtering Eqn. (1); enforce *locality* with polynomial filters Eqn. (2); specialize to a first-order approximation and absorb the identity with self-loops to obtain the *renormalized* propagator $\hat{\mathbf{A}}$ Eqn. (4); compose with a learnable feature mixer to yield the GCN layer Eqn. (5). The layer is permutation equivariant, k -hop localized, and implements neighborhood averaging (Laplacian smoothing). We provided full backpropagation Eqns. (9)–(11), training objective Eqn. (7), algorithmic steps, and variable dimensions for a complete, implementation-ready derivation.