

Recurrent Neural Network - Derivations & Proofs

Paul F. Roysdon, Ph.D.

Contents

1 Mathematical Derivations & Proofs	1
1.1 Introduction	1
1.2 Data and Notation	1
1.3 Model Formulation (Vanilla RNN)	2
1.4 Training Objective (Empirical Risk)	2
1.5 BPTT: Full Derivation	2
1.6 BPTT: Summary	4
1.7 Vanishing/Exploding Gradients: Spectral Analysis	5
1.8 Optimization and Practical Variants	5
1.9 Algorithm (Vanilla RNN + BPTT)	6
1.10 Computational Aspects	6
1.11 Summary of Variables and Their Dimensions	6
1.12 Summary	6

1 Mathematical Derivations & Proofs

1.1 Introduction

A **Recurrent Neural Network (RNN)** models sequential data by maintaining a latent *hidden state* that is updated recurrently as new inputs arrive. At each time step, the new hidden state is a nonlinear function of the current input and the previous hidden state, and predictions are produced from the hidden state. Training proceeds by minimizing a sequence loss via *Backpropagation Through Time* (BPTT), which computes exact gradients by unrolling the computation graph over time. We derive the forward equations, the full BPTT gradient recursions, and analyze vanishing/exploding gradients from first principles.

1.2 Data and Notation

Let a sequence be $\{(\mathbf{x}_t, \mathbf{y}_t)\}_{t=1}^T$, where

$$\mathbf{x}_t \in \mathbb{R}^{d_x} \text{ (input at time } t\text{)}, \quad \mathbf{y}_t \in \{0, 1\}^K \text{ (one-hot class target or a real vector).}$$

Hidden states $\mathbf{h}_t \in \mathbb{R}^{d_h}$. Parameters:

$$\mathbf{W}_{xh} \in \mathbb{R}^{d_h \times d_x}, \quad \mathbf{W}_{hh} \in \mathbb{R}^{d_h \times d_h}, \quad \mathbf{b}_h \in \mathbb{R}^{d_h}, \quad \mathbf{W}_{hy} \in \mathbb{R}^{K \times d_h}, \quad \mathbf{b}_y \in \mathbb{R}^K.$$

We denote the elementwise nonlinearity by $\phi(\cdot)$ (e.g., tanh or ReLU). The Hadamard (elementwise) product is \odot ; $\text{Diag}(\cdot)$ forms a diagonal matrix from a vector.

1.3 Model Formulation (Vanilla RNN)

Given an initial hidden \mathbf{h}_0 (zero or learned), the forward dynamics are

$$\text{Pre-activation: } \mathbf{s}_t = \mathbf{W}_{xh} \mathbf{x}_t + \mathbf{W}_{hh} \mathbf{h}_{t-1} + \mathbf{b}_h \in \mathbb{R}^{d_h}, \quad (1)$$

$$\text{Hidden update: } \mathbf{h}_t = \phi(\mathbf{s}_t) \in \mathbb{R}^{d_h}, \quad (2)$$

$$\text{Output logits: } \mathbf{z}_t = \mathbf{W}_{hy} \mathbf{h}_t + \mathbf{b}_y \in \mathbb{R}^K. \quad (3)$$

For classification, the predictive distribution at time t is the softmax

$$\hat{\mathbf{p}}_t = \text{softmax}(\mathbf{z}_t), \quad [\hat{\mathbf{p}}_t]_k = \frac{e^{z_{t,k}}}{\sum_{j=1}^K e^{z_{t,j}}}.$$

Sequence topologies. *Many-to-many* (per-time outputs): use all \mathbf{z}_t . *Many-to-one* (sequence classification): use only the final (or pooled) hidden \mathbf{h}_T for prediction via Eqn. (3) with $t = T$.

1.4 Training Objective (Empirical Risk)

For sequence classification with per-step cross-entropy, the loss of one sequence is

$$\mathcal{L} = \sum_{t=1}^T \ell_t, \quad \ell_t = -\mathbf{y}_t^\top \log \hat{\mathbf{p}}_t = -\sum_{k=1}^K y_{t,k} \log [\hat{\mathbf{p}}_t]_k. \quad (4)$$

(For many-to-one, drop the sum and keep only $t = T$.) Optionally add weight decay:

$$\frac{\lambda}{2} (\|\mathbf{W}_{xh}\|_F^2 + \|\mathbf{W}_{hh}\|_F^2 + \|\mathbf{W}_{hy}\|_F^2).$$

1.5 BPTT: Full Derivation

Loss Function

Suppose we have a loss function $\ell(\mathbf{y}_t, \hat{\mathbf{y}}_t)$ at each time step, where $\hat{\mathbf{y}}_t$ is the target output. The total loss over the sequence is given by

$$\mathcal{L} = \sum_{t=1}^T \ell(\mathbf{y}_t, \hat{\mathbf{y}}_t).$$

Our goal is to compute the gradients of \mathcal{L} with respect to the network parameters (e.g., \mathbf{W}_{hh} , \mathbf{W}_{xh} , and \mathbf{W}_{hy}).

Notation and Setup

For compactness, we define the pre-activation (or “net input”) at time t as

$$\mathbf{z}_t = \mathbf{W}_{hh} \mathbf{h}_{t-1} + \mathbf{W}_{xh} \mathbf{x}_t + \mathbf{b}_h,$$

so that

$$\mathbf{h}_t = \phi(\mathbf{z}_t).$$

Because the network is recurrent, \mathbf{h}_t influences not only the loss at time t (via \mathbf{y}_t) but also the losses at future time steps.

Gradient with Respect to \mathbf{W}_{hh}

The total gradient with respect to \mathbf{W}_{hh} is given by summing the contributions from each time step:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}}|_t.$$

Since \mathbf{W}_{hh} affects the hidden state at time t via

$$\mathbf{h}_t = \phi(\mathbf{z}_t) \quad \text{with} \quad \mathbf{z}_t = \mathbf{W}_{hh} \mathbf{h}_{t-1} + \dots,$$

by the chain rule we have

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}} = \sum_{t=1}^T \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{W}_{hh}}.$$

Define the “error” at the hidden state as

$$\delta_t^h \triangleq \frac{\partial \mathcal{L}}{\partial \mathbf{h}_t}.$$

Because \mathbf{h}_t contributes directly to the loss at time t and indirectly to future losses, we obtain the recursive relation:

$$\delta_t^h = \underbrace{\frac{\partial \ell_t}{\partial \mathbf{h}_t}}_{\text{direct contribution}} + \underbrace{\left(\frac{\partial \mathcal{L}}{\partial \mathbf{h}_{t+1}} \frac{\partial \mathbf{h}_{t+1}}{\partial \mathbf{h}_t} \right)}_{\text{future contributions}}.$$

That is, for $t = 1, \dots, T-1$,

$$\delta_t^h = \frac{\partial \ell_t}{\partial \mathbf{h}_t} + \delta_{t+1}^h \frac{\partial \mathbf{h}_{t+1}}{\partial \mathbf{h}_t}.$$

For the final time step $t = T$, we have:

$$\delta_T^h = \frac{\partial \ell_T}{\partial \mathbf{h}_T}.$$

Next, since

$$\mathbf{h}_t = \phi(\mathbf{z}_t) \quad \text{with} \quad \mathbf{z}_t = \mathbf{W}_{hh} \mathbf{h}_{t-1} + \dots,$$

the chain rule gives:

$$\frac{\partial \mathbf{h}_t}{\partial \mathbf{z}_t} = \phi'(\mathbf{z}_t).$$

(When ϕ is applied elementwise, $\phi'(\mathbf{z}_t)$ is a diagonal matrix with the derivatives on the diagonal.) Also, because

$$\mathbf{z}_t = \mathbf{W}_{hh} \mathbf{h}_{t-1} + \dots,$$

we have:

$$\frac{\partial \mathbf{z}_t}{\partial \mathbf{W}_{hh}} = \mathbf{h}_{t-1}^T.$$

Thus, the gradient of \mathbf{h}_t with respect to \mathbf{W}_{hh} is:

$$\frac{\partial \mathbf{h}_t}{\partial \mathbf{W}_{hh}} = \phi'(\mathbf{z}_t) \mathbf{h}_{t-1}^T.$$

Hence, the contribution at time t is:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}}|_t = \delta_t^h \phi'(\mathbf{z}_t) \mathbf{h}_{t-1}^T.$$

Thus, the overall gradient is:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}} = \sum_{t=1}^T \delta_t^h \phi'(\mathbf{z}_t) \mathbf{h}_{t-1}^T.$$

Unrolling the Error Back in Time

Because the hidden state at time t influences not only the loss at time t but also losses at later times, the total error δ_t^h can be written explicitly in terms of the losses at all future time steps. Unrolling the recursion yields:

$$\delta_t^h = \frac{\partial \ell_t}{\partial \mathbf{h}_t} + \frac{\partial \ell_{t+1}}{\partial \mathbf{h}_{t+1}} \frac{\partial \mathbf{h}_{t+1}}{\partial \mathbf{h}_t} + \frac{\partial \ell_{t+2}}{\partial \mathbf{h}_{t+2}} \frac{\partial \mathbf{h}_{t+2}}{\partial \mathbf{h}_t} + \dots,$$

or, more compactly,

$$\delta_t^h = \sum_{k=t}^T \frac{\partial \ell_k}{\partial \mathbf{h}_k} \left(\prod_{j=t+1}^k \frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}} \right).$$

Since each Jacobian $\frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}}$ is given by

$$\frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}} = \phi'(\mathbf{z}_j) \mathbf{W}_{hh},$$

we have:

$$\prod_{j=t+1}^k \frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}} = \prod_{j=t+1}^k [\phi'(\mathbf{z}_j) \mathbf{W}_{hh}],$$

with the convention that when $k = t$, the product is the identity.

Thus, the full expression for the gradient of the loss with respect to \mathbf{W}_{hh} becomes:

$$\boxed{\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}} = \sum_{t=1}^T \left[\left(\sum_{k=t}^T \frac{\partial \ell_k}{\partial \mathbf{h}_k} \prod_{j=t+1}^k [\phi'(\mathbf{z}_j) \mathbf{W}_{hh}] \right) \phi'(\mathbf{z}_t) \mathbf{h}_{t-1}^T \right].}$$

A similar derivation applies for the gradients with respect to \mathbf{W}_{xh} and the bias \mathbf{b}_h .

1.6 BPTT: Summary

We unroll Eqns. (1)–(3) over $t = 1:T$ and apply reverse-mode differentiation.

Local derivatives. Define the output error at time t (softmax with cross-entropy):

$$\delta_t^{(z)} \triangleq \frac{\partial \ell_t}{\partial \mathbf{z}_t} = \hat{\mathbf{p}}_t - \mathbf{y}_t \in \mathbb{R}^K. \quad (5)$$

Propagating to hidden \mathbf{h}_t through the linear head:

$$\mathbf{g}_t \triangleq \frac{\partial \ell}{\partial \mathbf{h}_t} = \mathbf{W}_{hy}^\top \delta_t^{(z)} + \underbrace{\left(\frac{\partial \ell}{\partial \mathbf{h}_t} \right)_{\text{from future}}}_{\text{to be accumulated}}. \quad (6)$$

Through the nonlinearity at time t :

$$\delta_t^{(s)} \triangleq \frac{\partial \ell}{\partial \mathbf{s}_t} = \mathbf{g}_t \odot \phi'(\mathbf{s}_t) \in \mathbb{R}^{d_h}. \quad (7)$$

Temporal recursion (error accumulation). Because \mathbf{h}_t influences *all* future losses via $\mathbf{h}_{t+1}, \mathbf{h}_{t+2}, \dots$ through \mathbf{W}_{hh} , the gradient $\left(\frac{\partial \ell}{\partial \mathbf{h}_t} \right)_{\text{from future}}$ satisfies the backward recurrence

$$\left(\frac{\partial \ell}{\partial \mathbf{h}_t} \right)_{\text{from future}} = \mathbf{W}_{hh}^\top \delta_{t+1}^{(s)} \quad (\text{with } \delta_{T+1}^{(s)} = \mathbf{0}). \quad (8)$$

Combining Eqns. (6)–(8), running $t = T, T-1, \dots, 1$,

$$\mathbf{g}_t = \mathbf{W}_{hy}^\top \delta_t^{(z)} + \mathbf{W}_{hh}^\top \delta_{t+1}^{(s)}, \quad \delta_t^{(s)} = \mathbf{g}_t \odot \phi'(\mathbf{s}_t). \quad (9)$$

Gradients w.r.t. parameters. Summing contributions over time gives

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hy}} = \sum_{t=1}^T \boldsymbol{\delta}_t^{(z)} \mathbf{h}_t^\top, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{b}_y} = \sum_{t=1}^T \boldsymbol{\delta}_t^{(z)}, \quad (10)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_{xh}} = \sum_{t=1}^T \boldsymbol{\delta}_t^{(s)} \mathbf{x}_t^\top, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{W}_{hh}} = \sum_{t=1}^T \boldsymbol{\delta}_t^{(s)} \mathbf{h}_{t-1}^\top, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{b}_h} = \sum_{t=1}^T \boldsymbol{\delta}_t^{(s)}. \quad (11)$$

The gradient w.r.t. the initial hidden (useful when it is learned) is $\frac{\partial \mathcal{L}}{\partial \mathbf{h}_0} = \mathbf{W}_{hh}^\top \boldsymbol{\delta}_1^{(s)}$.

Proof (chain rule over time). Differentiate $\ell = \sum_t \ell_t(\mathbf{z}_t(\mathbf{h}_t), \dots)$; apply chain rule

$$\frac{\partial \ell}{\partial \mathbf{h}_t} = \frac{\partial \ell_t}{\partial \mathbf{h}_t} + \left(\frac{\partial \mathbf{h}_{t+1}}{\partial \mathbf{h}_t} \right)^\top \frac{\partial \ell}{\partial \mathbf{h}_{t+1}} = \mathbf{W}_{hy}^\top \boldsymbol{\delta}_t^{(z)} + (\text{Diag}(\phi'(\mathbf{s}_{t+1})) \mathbf{W}_{hh})^\top \boldsymbol{\delta}_{t+1}^{(s)},$$

which yields Eqn. (9); parameter gradients follow from local Jacobians $\partial \mathbf{s}_t / \partial \mathbf{W}_{xh} = (\cdot) \mathbf{x}_t^\top$, $\partial \mathbf{s}_t / \partial \mathbf{W}_{hh} = (\cdot) \mathbf{h}_{t-1}^\top$, etc. ■

1.7 Vanishing/Exploding Gradients: Spectral Analysis

Consider the linearized recurrence around a trajectory:

$$\frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_{t-1}} = \text{Diag}(\phi'(\mathbf{s}_t)) \mathbf{W}_{hh} \triangleq \mathbf{J}_t.$$

The backprop factor from time $t+k$ to t is the product $\mathbf{J}_{t+k}^\top \cdots \mathbf{J}_{t+1}^\top$. Hence, for any consistent matrix norm,

$$\left\| \frac{\partial \ell}{\partial \mathbf{h}_t} \right\| \leq \sum_{k=0}^{T-t} \|\mathbf{W}_{hy}^\top\| \|\mathbf{J}_{t+1}^\top \cdots \mathbf{J}_{t+k}^\top\| \|\boldsymbol{\delta}_{t+k}^{(z)}\|.$$

If $\|\mathbf{J}_\tau\| \leq \rho < 1$ uniformly (e.g., $\|\mathbf{W}_{hh}\| \max \|\phi'\| < 1$), products decay $\mathcal{O}(\rho^k)$ (*vanishing gradients*). If $\rho > 1$, they grow (*exploding gradients*). In the linear RNN with $\phi(\cdot) = \text{id}$, $\mathbf{J}_t = \mathbf{W}_{hh}$ and $\partial \ell / \partial \mathbf{h}_t$ contains $(\mathbf{W}_{hh}^\top)^k$; behavior is governed by the spectral radius $\rho(\mathbf{W}_{hh})$.

Mitigations. Gradient clipping (e.g., rescale when $\|\nabla\|_2 > c$), orthogonal/identity initialization of \mathbf{W}_{hh} , gated architectures (LSTM/GRU) that modify \mathbf{J}_t to keep eigenvalues near 1, careful choice of ϕ (e.g., tanh vs. ReLU) and normalization.

1.8 Optimization and Practical Variants

- **Truncated BPTT.** For long sequences, accumulate gradients over windows of length $L \ll T$ by periodically detaching the computation graph: preserve local dependencies while controlling memory and compute.
- **Teacher forcing / scheduled sampling.** For auto-regressive outputs, during training feed ground-truth past outputs; at inference, feed model outputs.
- **Regularization.** Weight decay, dropout on inputs/hidden states, gradient clipping, early stopping.
- **Initialization.** \mathbf{W}_{hh} orthogonal or scaled identity; \mathbf{b}_h sometimes positive for ReLU to avoid dead states; \mathbf{h}_0 zeros or learned.

1.9 Algorithm (Vanilla RNN + BPTT)

1. **Input:** sequence $\{(\mathbf{x}_t, \mathbf{y}_t)\}_{t=1}^T$, learning rate η , (optional) truncation length L .
2. **Forward:** set \mathbf{h}_0 ; for $t = 1:T$ compute Eqns. (1)–(3), $\hat{\mathbf{p}}_t$, and accumulate \mathcal{L} via Eqn. (4).
3. **Backward:** initialize $\boldsymbol{\delta}_{T+1}^{(s)} = \mathbf{0}$. For $t = T:1$:
 - (a) $\boldsymbol{\delta}_t^{(z)} = \hat{\mathbf{p}}_t - \mathbf{y}_t$ (Eqn. (5))
 - (b) $\mathbf{g}_t = \mathbf{W}_{hy}^\top \boldsymbol{\delta}_t^{(z)} + \mathbf{W}_{hh}^\top \boldsymbol{\delta}_{t+1}^{(s)}$; $\boldsymbol{\delta}_t^{(s)} = \mathbf{g}_t \odot \phi'(\mathbf{s}_t)$ (Eqn. (9))
 - (c) Accumulate gradients using Eqns. (10)–(11).
4. **Update:** $\Theta \leftarrow \Theta - \eta \nabla \mathcal{L}$ with SGD/Adam (with clipping if needed).

1.10 Computational Aspects

Per-step costs: forward $\mathcal{O}(d_h d_x + d_h^2 + K d_h)$; backward comparable. Memory $\mathcal{O}(Td_h)$ to store $\{\mathbf{s}_t, \mathbf{h}_t\}$ unless using checkpointing or truncation.

1.11 Summary of Variables and Their Dimensions

- $\mathbf{x}_t \in \mathbb{R}^{d_x}$: input at time t ; $\mathbf{y}_t \in \{0, 1\}^K$ (one-hot) or \mathbb{R}^K (real target).
- $\mathbf{h}_t \in \mathbb{R}^{d_h}$: hidden state; $\mathbf{s}_t \in \mathbb{R}^{d_h}$: pre-activation; $\mathbf{z}_t \in \mathbb{R}^K$: logits.
- $\hat{\mathbf{p}}_t \in \mathbb{R}^K$: softmax probabilities.
- $\mathbf{W}_{xh} \in \mathbb{R}^{d_h \times d_x}$, $\mathbf{W}_{hh} \in \mathbb{R}^{d_h \times d_h}$, $\mathbf{b}_h \in \mathbb{R}^{d_h}$.
- $\mathbf{W}_{hy} \in \mathbb{R}^{K \times d_h}$, $\mathbf{b}_y \in \mathbb{R}^K$.
- Backprop errors: $\boldsymbol{\delta}_t^{(z)} \in \mathbb{R}^K$, $\mathbf{g}_t \in \mathbb{R}^{d_h}$, $\boldsymbol{\delta}_t^{(s)} \in \mathbb{R}^{d_h}$.

1.12 Summary

From first principles, a vanilla RNN iterates the nonlinear state update $\mathbf{h}_t = \phi(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h)$ and produces logits $\mathbf{z}_t = \mathbf{W}_{hy}\mathbf{h}_t + \mathbf{b}_y$. Unrolling over time and applying the chain rule yields exact BPTT: the hidden-state adjoint obeys the backward recursion $\mathbf{g}_t = \mathbf{W}_{hy}^\top(\hat{\mathbf{p}}_t - \mathbf{y}_t) + \mathbf{W}_{hh}^\top \boldsymbol{\delta}_{t+1}^{(s)}$, with $\boldsymbol{\delta}_t^{(s)} = \mathbf{g}_t \odot \phi'(\mathbf{s}_t)$, and the parameter gradients are simple outer-product accumulations over time. A spectral analysis of the Jacobian product explains vanishing/exploding gradients and motivates practical remedies (clipping, orthogonal init, gating). This provides a complete, dimensionally explicit foundation for implementing and training vanilla RNNs with BPTT.